

Workstream Comparison: Alcott Food Deserts (GPT-5 vs GPT-5-nano)

Sources

- Baseline (archive): [reports/archived/workflow-report-d1617aa16048423daa22ca7e29c3526c.json](#)
- Variant (current): [reports/workflow-report-e53e1c45434547afaedf7a0d079d5615.json](#)

Workstream Table Schema

| Column | Type | Purpose |
|-----------------------|----------------------------|--|
| comparison_group_id | text | Logical key that groups runs into one composed workstream component. |
| source_bucket | 'enum(archive current)' | |
| report_path | text | Path to the source workflow-report JSON used for this row. |
| run_id | text | Unique workflow execution identifier. |
| papers_dir | text | Input paper path(s) processed by the run. |
| started_at | timestamptz | Workflow start timestamp. |
| finished_at | timestamptz | Workflow end timestamp. |
| duration_sec | numeric | Run duration in seconds for cost/latency comparison. |
| chat_model | text | Primary LLM used by the run. |
| embedding_model | text | Embedding model used for retrieval/chunk similarity. |
| report_question_set | enum | Question set mode (structured) |
| report_question_count | integer | Number of structured/agentic report question records emitted. |
| agentic_status | text | Top-level status for agentic stage. |
| index_status | text | Top-level status for index stage. |
| index_reason_or_error | text | Failure/skipped reason to diagnose differences. |
| report_store_status | text | Whether report was stored to Postgres. |
| prep_corpus_hash | text | Fingerprint of corpus content; validates same input set. |
| config_hash | text | Effective configuration hash; validates run comparability. |
| confidence_mean | numeric | Mean confidence score across structured questions. |
| confidence_labels | text | Label distribution summary for quick quality profile. |
| final_answer_preview | text | Truncated synthesis preview for side-by-side semantic comparison. |

Composed Workstream Rows

| comparison_group_id | source_bucket | report_path |
|--|---------------|--|
| alcott-food-deserts-gpt5-vs-gpt5-nano-v1 | archive | reports/archived/workflow-report-d1617aa16048423daa22ca7e29c3526c.json |
| alcott-food-deserts-gpt5-vs-gpt5-nano-v1 | current | reports/workflow-report-e53e1c45434547afaedf7a0d079d5615.json |

Focused Comparison (Model)

| Metric | GPT-5 (baseline) | GPT-5-nano (variant) |
|-----------------------|----------------------------------|----------------------------------|
| run_id | d1617aa16048423daa22ca7e29c3526c | e53e1c45434547afaedf7a0d079d5615 |
| chat_model | gpt-5 | gpt-5-nano |
| duration_sec | 845.17 | 406.22 |
| report_question_count | 84 | 84 |
| confidence_mean | 0.3041450240118245 | 0.3042301084348596 |
| confidence_labels | high=0, medium=18, low=65 | high=0, medium=18, low=65 |
| index_status | skipped | failed |
| report_store_status | skipped | stored |

Notes

- Config hash differs between runs (expected when model defaults changed).
- GPT-5-nano run completed full agentic+structured flow (`report_question_set=both`).
- New nano audit report: `reports/audit-workflow-report-e53e1c45434547afaedf7a0d079d5615.md`.